


Alzheimer's MRI Classification — Figures

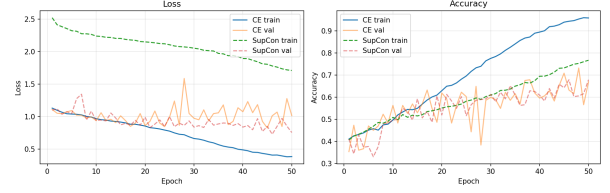
ResNet-10 vs ResNet-18

Contents

1 Training Curves



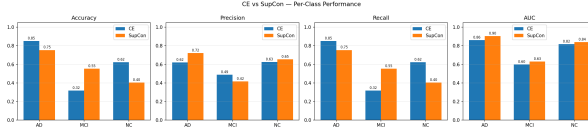
(a) ResNet-10



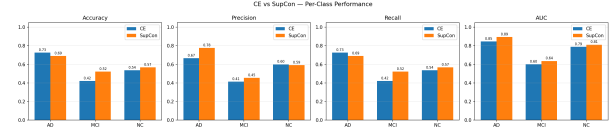
(b) ResNet-18

Figure 1: Training and validation loss/accuracy curves for CE and SupCon runs.

2 CE vs. SupCon Comparison



(a) ResNet-10



(b) ResNet-18

Figure 2: Side-by-side metric comparison between cross-entropy (CE) and supervised contrastive (SupCon) training.

3 Confusion Matrices

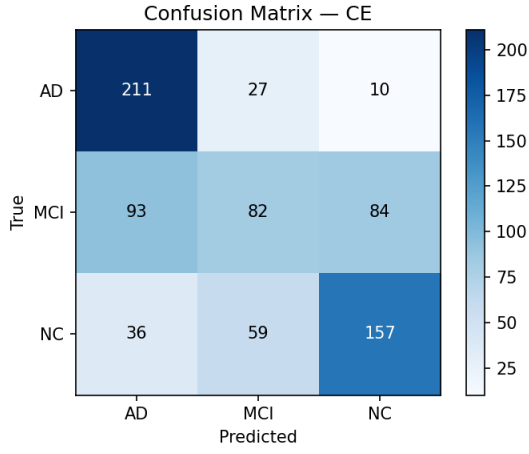
3.1 Cross-Entropy

3.2 Supervised Contrastive

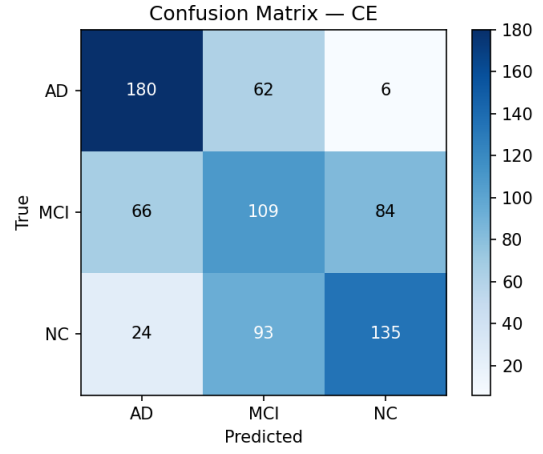
4 ROC Curves

4.1 Cross-Entropy

4.2 Supervised Contrastive

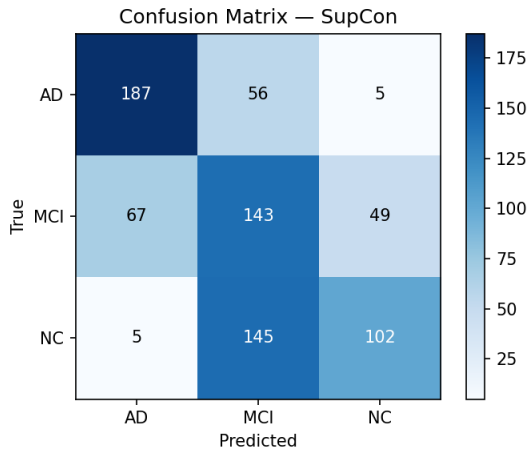


(a) ResNet-10 CE

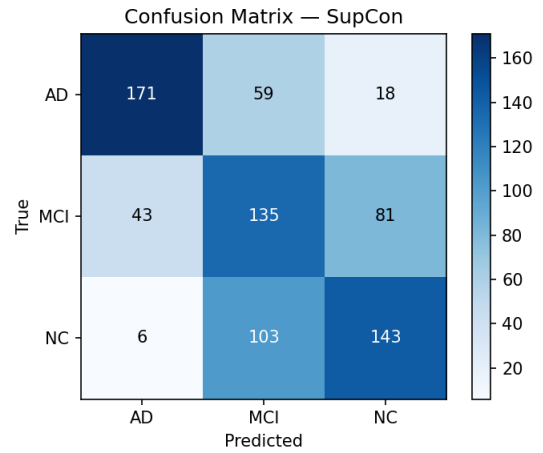


(b) ResNet-18 CE

Figure 3: Confusion matrices for the cross-entropy baseline.

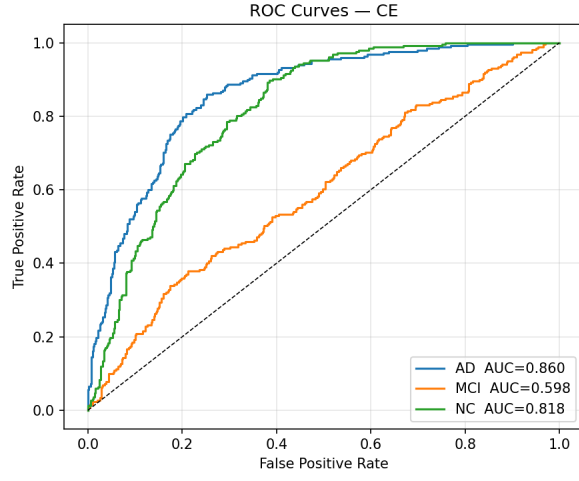


(a) ResNet-10 SupCon

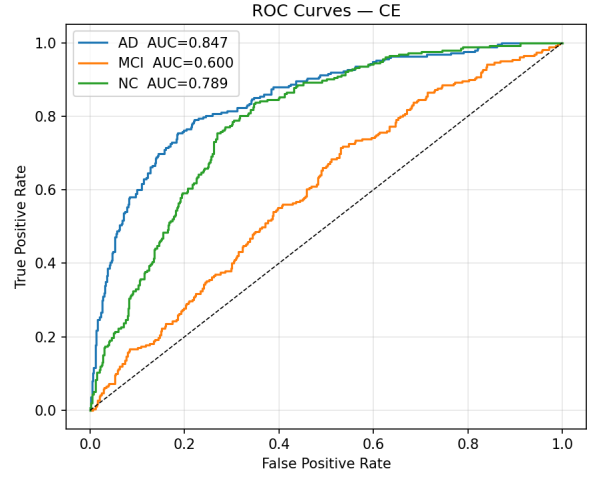


(b) ResNet-18 SupCon

Figure 4: Confusion matrices for supervised contrastive training.

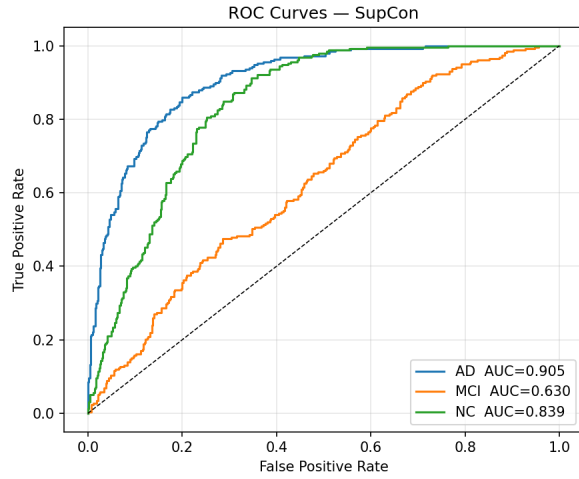


(a) ResNet-10 CE

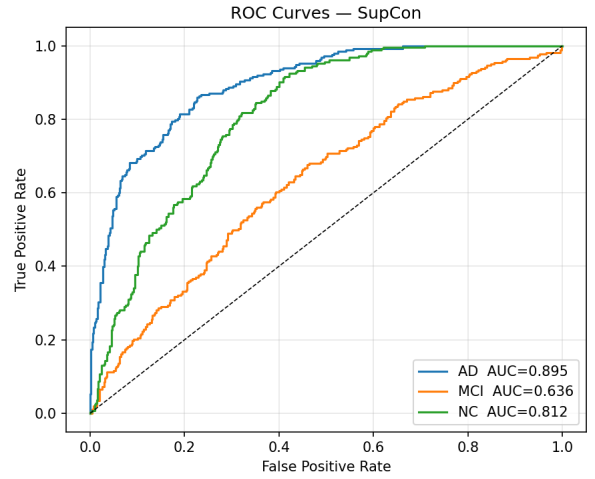


(b) ResNet-18 CE

Figure 5: ROC curves (one-vs-rest) for the cross-entropy baseline.



(a) ResNet-10 SupCon



(b) ResNet-18 SupCon

Figure 6: ROC curves (one-vs-rest) for supervised contrastive training.